

CS5613

Neural Networks

Assignment 2 - CNN Group Assignment

Group Name : 199342EC.G.Madage-199355VW.Perera-199361KR.Senanayaka

199342E - C.G. Madage

199355V - Wishwa Perera

199361K- M.D.R.N. Senanayake

MSc in Computer Science 2019

Department of Computer Science and Engineering

University of Moratuwa

Q1) Answer for Question1 is published in the [Github](https://github.com/randikacse/Assignment_CNN_Group/blob/main/question01.ipynb)

https://github.com/randikacse/Assignment_CNN_Group/blob/main/question01.ipynb

randikacse / Assignment_CNN_Group

Unwatch 1 Star 0 Fork 0

<> Code Issues Pull requests Actions Projects Wiki Security Insights Settings

main Assignment_CNN_Group / question01.ipynb

Go to file

randikacse Plant Seedling Classification using Convolutional Neural Networks

Latest commit 605c7bb 8 minutes ago History

1 contributor

1 lines (1 sloc) 613 KB

<> Raw Blame

Plant Seedling Classification using Convolutional Neural Networks

First need to add input data, for that click on 'Add Data'. Go to the competitions tab and then search for the 'Plant Seedlings Classification' dataset and add it. This dataset contains a training set and a testing set of images of plant seedlings at various growth stages. Each image has its own unique ID. The dataset has 12 main plant species which we need to classify the testing set into. For this task, first, we processed and cleaned the data using image processing techniques. Then we build a model and evaluate it.

Import Libraries

```
In [1]: import cv2
from glob import glob
import numpy as np
from matplotlib import pyplot as plt
import pandas as pd
```

Import training data set

```
In [2]: train_file_path = '/kaggle/input/plant-seedlings-classification/train/*/*.png'
train_images = glob(train_file_path)
```

Preprocess the training images set

```
In [3]: scale_to = 100 # pixel of image to re-scale
num_of_seed = 7 # fixing random seeds

train_img_array = []
train_label_array = []
count = 1
num = len(train_images)

# Image resizing, get all labels
for img in train_images:
    print(str(count) + "/" + str(num), end="\n")
    train_img_array.append(cv2.resize(cv2.imread(img), (scale_to, scale_to))) # Get image (with resizing)
    img_array = img.split('/')
    train_label_array.append(img_array[5]) # image type
    count += 1

train_images = np.asarray(train_img_array) # Train images set
train_labels = pd.DataFrame(train_label_array) # Train Labels set

4750/4750

Print processed images
```

```
In [4]: for count in range(8):
plt.subplot(2, 4, count + 1)
plt.imshow(train_images[count])
```



Clean the training images set

```
In [5]: cleaned_train_images = []
show_samples = True
for img in train_images:
    # gaussian blur
    blur_img = cv2.GaussianBlur(img, (5, 5), 0)

    # convert to HSV image
    hsvImg = cv2.cvtColor(blur_img, cv2.COLOR_BGR2HSV)

    # Create mask (parameters - green color range)
```

Clean the training images set

```
In [5]: cleaned_train_images = []
show_samples = True
for img in train_images:
    # gaussian blur
    blur_img = cv2.GaussianBlur(img, (5, 5), 0)

    # convert to HSV image
    hsvImg = cv2.cvtColor(blur_img, cv2.COLOR_BGR2HSV)

    # Create mask (parameters - green color range)
    lower_green = (25, 41, 50)
    upper_green = (74, 255, 255)
    mask = cv2.inRange(hsvImg, lower_green, upper_green)
    kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (11, 11))
    mask = cv2.morphologyEx(mask, cv2.MORPH_CLOSE, kernel)


    # Create bool mask
    bMask = mask > 0

    # Apply the mask
    cleaned = np.zeros_like(img, np.uint8) # Create empty image
    cleaned[bMask] = img[bMask] # Apply boolean mask to the origin image

    cleaned_train_images.append(cleaned) # Append image without background


# Show examples
if show_samples:
    plt.subplot(2, 3, 1); plt.imshow(img) # Show the original image
    plt.subplot(2, 3, 2); plt.imshow(blur_img) # Blur image
    plt.subplot(2, 3, 3); plt.imshow(hsvImg) # HSV image
    plt.subplot(2, 3, 4); plt.imshow(mask) # Mask
    plt.subplot(2, 3, 5); plt.imshow(bMask) # Boolean mask
    plt.subplot(2, 3, 6); plt.imshow(cleaned) # Image without background
    show_samples = False

cleaned_train_img = np.asarray(cleaned_train_images)
```



Print cleaned images

```
In [6]: for i in range(8):
plt.subplot(2, 4, i + 1)
plt.imshow(cleaned_train_img[i])
```



```
In [7]: cleaned_train_imgs = cleaned_train_img / 255
```

Identify the Classes

```
In [8]: from keras.utils import np_utils
from sklearn import preprocessing
import matplotlib.pyplot as plt

# Encode labels and create classes
le = preprocessing.LabelEncoder()
le.fit(train_labels[0])
print("Classes: " + str(le.classes_))
encode_train_labels = le.transform(train_labels[0])


# Make labels categorical
cleaned_train_label = np_utils.to_categorical(encode_train_labels)
num_classes = cleaned_train_label.shape[1]
print("num of classes: " + str(num_classes))

# Plot of Label types numbers
train_labels[0].value_counts().plot(kind='bar')

Classes: ['Black-grass' 'Charlock' 'Cleavers' 'Common Chickweed' 'Common wheat'
'Fat Hen' 'Loose Silky-bent' 'Maize' 'Scentless Mayweed'
'Shepherds Purse' 'Small-flowered Cranesbill' 'Sugar beet']
num of classes: 12
```


Out[8]: <AxesSubplot:>

Out[8]: <AxesSubplot:>



```
In [9]: train_labels[0].value_counts().plot(kind='pie')
```

Out[9]: <AxesSubplot:ylabel='0'>



Split train and test data set

```
In [10]: from sklearn.model_selection import train_test_split

trainX, testX, trainY, testY = train_test_split(cleaned_train_imgs, cleaned_train_label,
                                                test_size=0.2, random_state=num_of_seed,
                                                stratify = cleaned_train_label)

In [11]: from keras.preprocessing.image import ImageDataGenerator
data_gen = ImageDataGenerator(
    rotation_range=180, # randomly rotate images in the range
    zoom_range = 0.1, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally
    height_shift_range=0.1, # randomly shift images vertically
    horizontal_flip=True, # randomly flip images horizontally
    vertical_flip=True # randomly flip images vertically
)
data_gen.fit(trainX)
```

Create the Model

```
In [12]: import numpy
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout
from keras.layers import Flatten
from keras.layers.convolutional import Conv2D
from keras.layers.convolutional import MaxPooling2D
from keras.layers import BatchNormalization

numpy.random.seed(num_of_seed) # num_of_seed

model = Sequential()

model.add(Conv2D(filters=64, kernel_size=(5, 5), input_shape=(scale_to, scale_to, 3), activation='relu'))
model.add(BatchNormalization(axis=3))
model.add(Conv2D(filters=64, kernel_size=(5, 5), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(BatchNormalization(axis=3))
model.add(Dropout(0.1))

model.add(Conv2D(filters=128, kernel_size=(5, 5), activation='relu'))
model.add(BatchNormalization(axis=3))
model.add(Conv2D(filters=128, kernel_size=(5, 5), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(BatchNormalization(axis=3))
model.add(Dropout(0.1))

model.add(Conv2D(filters=256, kernel_size=(5, 5), activation='relu'))
model.add(BatchNormalization(axis=3))
model.add(Conv2D(filters=256, kernel_size=(5, 5), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(BatchNormalization(axis=3))
model.add(Dropout(0.1))

model.add(Flatten())

model.add(Dense(256, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(256, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(num_classes, activation='softmax'))

model.summary()

# compile model
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 96, 96, 64)	4864
batch_normalization (BatchNo	(None, 96, 96, 64)	256
conv2d_1 (Conv2D)	(None, 92, 92, 64)	102464
max_pooling2d (MaxPooling2D)	(None, 46, 46, 64)	0
batch_normalization_1 (Batch	(None, 46, 46, 64)	256
dropout (Dropout)	(None, 46, 46, 64)	0
conv2d_2 (Conv2D)	(None, 42, 42, 128)	204928
batch_normalization_2 (Batch	(None, 42, 42, 128)	512
conv2d_3 (Conv2D)	(None, 38, 38, 128)	409728
max_pooling2d_1 (MaxPooling2	(None, 19, 19, 128)	0
batch_normalization_3 (Batch	(None, 19, 19, 128)	512
dropout_1 (Dropout)	(None, 19, 19, 128)	0
conv2d_4 (Conv2D)	(None, 15, 15, 256)	819456
batch_normalization_4 (Batch	(None, 15, 15, 256)	1024
conv2d_5 (Conv2D)	(None, 11, 11, 256)	1638656
max_pooling2d_2 (MaxPooling2	(None, 5, 5, 256)	0
batch_normalization_5 (Batch	(None, 5, 5, 256)	1024
dropout_2 (Dropout)	(None, 5, 5, 256)	0
flatten (Flatten)	(None, 6400)	0
dense (Dense)	(None, 256)	1638656
batch_normalization_6 (Batch	(None, 256)	1024
dropout_3 (Dropout)	(None, 256)	0

```
In [13]: from keras.callbacks import ModelCheckpoint, ReduceLRonPlateau, CSVLogger

# Learning rate reduction
learning_rate_reduction = ReduceLRonPlateau(monitor='val_accuracy',
                                             patience=3,
                                             verbose=1,
                                             factor=0.4,
                                             min_lr=0.00001)

# add check points
file_path = "/kaggle/working/weights.best_{epoch:02d}-{val_accuracy:.2f}.hdf5"

check_point = ModelCheckpoint(file_path, monitor='val_accuracy',
                              verbose=1, save_best_only=True, mode='max')

file_path = "/kaggle/working/weights.last_auto.hdf5"
checkpoint_all = ModelCheckpoint(file_path, monitor='val_accuracy',
                                verbose=1, save_best_only=False, mode='max')

# all callbacks
callbacks_list = [check_point, learning_rate_reduction, checkpoint_all]

# fit model
hist = model.fit_generator(data_gen.flow(trainX, trainY, batch_size=75), epochs=30, validation_data=(testX, testY), callbacks=callbacks_list)

/opt/conda/lib/python3.7/site-packages/tensorflow/python/keras/engine/training.py:1844: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
  warnings.warn("`Model.fit_generator` is deprecated and ")

Epoch 1/30
51/51 [=====] - 75s 1s/step - loss: 3.0351 - accuracy: 0.2073 - val_loss: 25.1081 - val_accuracy: 0.0463

Epoch 00001: val_accuracy improved from -inf to 0.04632, saving model to /kaggle/working/weights.best_01-0.05.hdf5

Epoch 00001: saving model to /kaggle/working/weights.last_auto.hdf5
Epoch 2/30
51/51 [=====] - 57s 1s/step - loss: 2.0388 - accuracy: 0.3809 - val_loss: 23.2107 - val_accuracy: 0.0463

Epoch 00002: val_accuracy did not improve from 0.04632

Epoch 00002: saving model to /kaggle/working/weights.last_auto.hdf5
Epoch 3/30
51/51 [=====] - 57s 1s/step - loss: 1.5278 - accuracy: 0.4813 - val_loss: 25.1781 - val_accuracy: 0.0611

Epoch 00003: val_accuracy improved from 0.04632 to 0.06105, saving model to /kaggle/working/weights.best_03-0.06.hdf5

Epoch 00003: saving model to /kaggle/working/weights.last_auto.hdf5
Epoch 4/30
51/51 [=====] - 57s 1s/step - loss: 1.4965 - accuracy: 0.4924 - val_loss: 27.5217 - val_accuracy: 0.0611

Epoch 00004: val_accuracy did not improve from 0.06105

Epoch 00004: saving model to /kaggle/working/weights.last_auto.hdf5
Epoch 5/30
51/51 [=====] - 57s 1s/step - loss: 1.2689 - accuracy: 0.5660 - val_loss: 26.1744 - val_accuracy: 0.0611

Epoch 00005: val_accuracy did not improve from 0.06105

Epoch 00005: saving model to /kaggle/working/weights.last_auto.hdf5
Epoch 6/30
51/51 [=====] - 57s 1s/step - loss: 1.1779 - accuracy: 0.5909 - val_loss: 33.1222 - val_accuracy: 0.0611

Epoch 00006: val_accuracy did not improve from 0.06105

Epoch 00006: ReduceLRonPlateau reducing learning rate to 0.0004000000189989805.

Epoch 00006: saving model to /kaggle/working/weights.last_auto.hdf5
Epoch 7/30
51/51 [=====] - 57s 1s/step - loss: 1.0414 - accuracy: 0.6396 - val_loss: 29.3030 - val_accuracy: 0.0611

Epoch 00007: val_accuracy did not improve from 0.06105

Epoch 00007: saving model to /kaggle/working/weights.last_auto.hdf5
Epoch 8/30
51/51 [=====] - 57s 1s/step - loss: 1.0099 - accuracy: 0.6626 - val_loss: 22.7306 - val_accuracy: 0.0611

Epoch 00008: val_accuracy did not improve from 0.06105

Epoch 00008: saving model to /kaggle/working/weights.last_auto.hdf5
Epoch 9/30
51/51 [=====] - 57s 1s/step - loss: 0.9299 - accuracy: 0.6829 - val_loss: 18.5631 - val_accuracy: 0.0611

Epoch 00009: val_accuracy did not improve from 0.06105

Epoch 00009: ReduceLRonPlateau reducing learning rate to 0.00016000000759959222.

Epoch 00009: saving model to /kaggle/working/weights.last_auto.hdf5
Epoch 10/30
51/51 [=====] - 57s 1s/step - loss: 0.8397 - accuracy: 0.7045 - val_loss: 16.0639 - val_accuracy: 0.0621
```

```
In [14]: model.load_weights("/kaggle/working/weights.last_auto.hdf5")
```

```
In [15]: print(model.evaluate(trainX, trainY)) # evaluate on train set
print(model.evaluate(testX, testY)) # evaluate on test set
```

```
119/119 [=====] - 45s 379ms/step - loss: 0.2981 - accuracy: 0.8921
[0.2981433868408203, 0.8921052813529968]
30/30 [=====] - 11s 362ms/step - loss: 1.8859 - accuracy: 0.8568
[1.8859095573425293, 0.8568421006202698]
```

```
In [16]: from sklearn.metrics import confusion_matrix
import itertools

def plot_confusion_matrix(cm, classes,
                           normalize = False,
                           title = 'Confusion matrix',
                           cmap = plt.cm.Blues):

    fig = plt.figure(figsize=(10,10))
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=90)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]


    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

# Predict the values from the validation dataset
pred_y = model.predict(testX)
pred_y_classes = np.argmax(pred_y, axis = 1)
true_y = np.argmax(testY, axis = 1)

# confusion matrix
confusion_MTX = confusion_matrix(true_y, pred_y_classes)

# plot the confusion matrix
plot_confusion_matrix(confusion_MTX, classes = le.classes_)
```



```
In [17]: test_images_path = '/kaggle/input/plant-seedlings-classification/test/*.png'
test_images = glob(test_images_path)
```

```

In [18]: test_img_array = []
test_id_array = []
count = 1
num = len(test_images)


# Obtain images and resizing, obtain Labels
for img in test_images:
    print("Obtain images: " + str(count) + "/" + str(num), end='\n')
    img_array = img.split('/')
    test_id_array.append(img_array[5]) # image id
    test_img_array.append(cv2.resize(cv2.imread(img), (scale_to, scale_to)))
    count += 1

test_imgs = np.asarray(test_img_array) # Train images set

for i in range(8):
    plt.subplot(2, 4, i + 1)
    plt.imshow(test_imgs[i])

```

Obtain images: 794/794




```

In [19]: cleaned_test_img = []
show_samples = True
for img in test_imgs:
    # gaussian blur
    blur_img = cv2.GaussianBlur(img, (5, 5), 0)

    # convert to HSV image
    hsvImg = cv2.cvtColor(blur_img, cv2.COLOR_BGR2HSV)

    # create mask (parameters - green color range)
    lower_green = (25, 41, 50)
    upper_green = (74, 255, 255)
    mask = cv2.inRange(hsvImg, lower_green, upper_green)
    kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (11, 11))
    mask = cv2.morphologyEx(mask, cv2.MORPH_CLOSE, kernel)

    # create bool mask
    bMask = mask > 0


    # Apply the mask
    cleaned = np.zeros_like(img, np.uint8) # Create empty image
    cleaned[bMask] = img[bMask] # Apply boolean mask to the origin image

    cleaned_test_img.append(cleaned) # Append image without background

# Show examples
if show_samples:
    plt.subplot(2, 3, 1); plt.imshow(img) # Show the original image
    plt.subplot(2, 3, 2); plt.imshow(blur_img) # Blur image
    plt.subplot(2, 3, 3); plt.imshow(hsvImg) # HSV image
    plt.subplot(2, 3, 4); plt.imshow(mask) # Mask
    plt.subplot(2, 3, 5); plt.imshow(bMask) # Boolean mask
    plt.subplot(2, 3, 6); plt.imshow(cleaned) # Image without background
    show_samples = False

cleaned_test_img = np.asarray(cleaned_test_img)


```



```

In [20]: for i in range(8):
plt.subplot(2, 4, i + 1)
plt.imshow(cleaned_test_img[i])

```



```

In [21]: cleaned_test_img = cleaned_test_img / 255

```

```

In [22]: pred = model.predict(cleaned_test_img)

```

Create submission

```

In [23]: predNum = np.argmax(pred, axis=1)
predStr = le.classes_[predNum]

res = {'file': test_id_array, 'species': predStr}
res = pd.DataFrame(res)

```

```

In [24]: res.to_csv("/kaggle/working/result_v2.csv", index=False)

```

Q2) Answer for Question2 is published in the [Github](https://github.com/randikacse/Assignment_CNN_Group/blob/main/question02.ipynb)

https://github.com/randikacse/Assignment_CNN_Group/blob/main/question02.ipynb

 randikacse Transfer learning with pretrained Keras model Xception 

Latest commit 81a66a8 25 minutes ago  History

1 contributor

1.16 MB   

Transfer learning with pretrained Keras model

In this notebook, we used a pre-trained Xception model to predict images in Plant Seedlings Classification competition.

In order to use the pre-trained weights for the Xception model, need to add a new dataset containing the weights. Go to the Data tab and click on 'Add Data'. Then search for the 'Keras Pretrained Model' dataset and add it. Also need to add input data, for that click on 'Add Data'. Go to the competitions tab and then search for the 'Plant Seedlings Classification' dataset and add it.

Importing all the necessary modules:

```
In [1]: %matplotlib inline
import matplotlib.pyplot as plt
plt.rcParams['figure.figsize'] = [16, 10]
plt.rcParams['font.size'] = 16
import numpy as np
import os
import pandas as pd
import seaborn as sns
from keras.applications import xception
from keras.preprocessing import image
from mpl_toolkits.axes_grid1 import ImageGrid
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
from tqdm import tqdm
```

```
In [2]: from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

Use Keras Pretrained Models dataset and copy the pretrained models(xception) to the cache directory (~/.keras/models)

```
In [3]: !ls ../input/keras-pretrained-models/

Kuzma.JPG
imagenet_class_index.json
inception_resnet_v2_weights_tf_dim_ordering_tf_kernels.h5
inception_resnet_v2_weights_tf_dim_ordering_tf_kernels_notop.h5
inception_v3_weights_tf_dim_ordering_tf_kernels.h5
inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5
resnet50_weights_tf_dim_ordering_tf_kernels.h5
resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5
vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
xception_weights_tf_dim_ordering_tf_kernels.h5
xception_weights_tf_dim_ordering_tf_kernels_notop.h5
```

```
In [4]: cache_dir = os.path.expanduser(os.path.join('~', '.keras'))
if not os.path.exists(cache_dir):
    os.makedirs(cache_dir)
models_dir = os.path.join(cache_dir, 'models')
if not os.path.exists(models_dir):
    os.makedirs(models_dir)
```

```
In [5]: !cp ../input/keras-pretrained-models/xception* ~/.keras/models/
```

```
In [6]: !ls ~/.keras/models

xception_weights_tf_dim_ordering_tf_kernels.h5
xception_weights_tf_dim_ordering_tf_kernels_notop.h5
```

Check the plant seedlings Dataset

```
In [7]: !ls ../input/plant-seedlings-classification

sample_submission.csv  test  train
```

```
In [8]: CATEGORIES = ['Black-grass', 'Charlock', 'Cleavers', 'Common Chickweed', 'Common wheat', 'Fat Hen', 'Loose Silky-bent',
                    'Maize', 'Scentless Mayweed', 'Shepherds Purse', 'Small-flowered Cranesbill', 'Sugar beet']
NUM_CATEGORIES = len(CATEGORIES)
```

```
In [9]: SAMPLE_PER_CATEGORY = 200
SEED = 1987
data_dir = '../input/plant-seedlings-classification/'
train_dir = os.path.join(data_dir, 'train')
test_dir = os.path.join(data_dir, 'test')
sample_submission = pd.read_csv(os.path.join(data_dir, 'sample_submission.csv'))
```

```
In [10]: sample_submission.head(2)
```

```
Out[10]:
```

	file	species
0	0021e90e4.png	Sugar beet
1	003d61042.png	Sugar beet

```
In [11]: for category in CATEGORIES:
          print('{} {} images'.format(category, len(os.listdir(os.path.join(train_dir, category)))))
```

```
Black-grass 263 images
Charlock 390 images
Cleavers 287 images
Common Chickweed 611 images
Common wheat 221 images
Fat Hen 475 images
Loose Silky-bent 654 images
Maize 221 images
Scentless Mayweed 516 images
Shepherds Purse 231 images
Small-flowered Cranesbill 496 images
Sugar beet 385 images
```

```
In [12]: train = []
          for category_id, category in enumerate(CATEGORIES):
              for file in os.listdir(os.path.join(train_dir, category)):
                  train.append(['train/{}/{}'.format(category, file), category_id, category])
          train = pd.DataFrame(train, columns=['file', 'category_id', 'category'])
          train.head(2)
          train.shape
```

```
Out[12]:
```

	file	category_id	category
0	train/Black-grass/2aa60045d.png	0	Black-grass
1	train/Black-grass/a47cfeec4.png	0	Black-grass

```
Out[12]: (4750, 3)
```

Training sample

```
In [13]: train = pd.concat([train[train['category'] == c][:SAMPLE_PER_CATEGORY] for c in CATEGORIES])
          train = train.sample(frac=1)
          train.index = np.arange(len(train))
          train.head(2)
          train.shape
```

```
Out[13]:
```

	file	category_id	category
0	train/Shepherds Purse/8e1efae9e.png	9	Shepherds Purse
1	train/Shepherds Purse/150ab985f.png	9	Shepherds Purse

```
Out[13]: (2400, 3)
```

```
In [14]: test = []
          for file in os.listdir(test_dir):
              test.append(['test/{}'.format(file), file])
          test = pd.DataFrame(test, columns=['filepath', 'file'])
          test.head(2)
          test.shape
```

```
Out[14]:
```

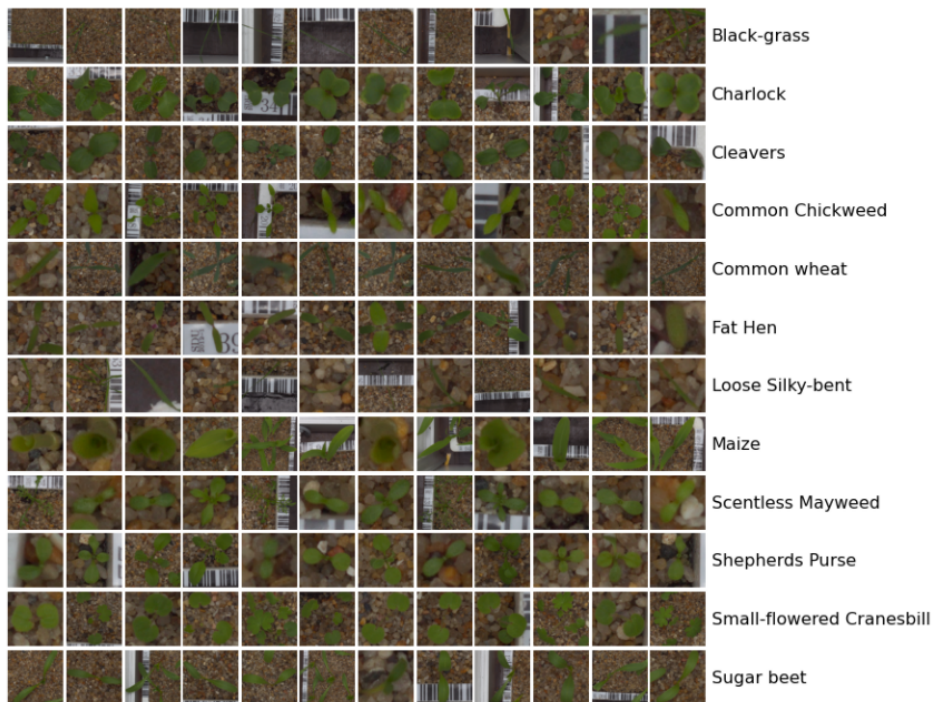
	filepath	file
0	test/fd87b36ae.png	fd87b36ae.png
1	test/0e8492cb1.png	0e8492cb1.png

```
Out[14]: (794, 2)
```

```
In [15]: def read_img(filepath, size):
          img = image.load_img(os.path.join(data_dir, filepath), target_size=size)
          img = image.img_to_array(img)
          return img
```

Example images

```
In [16]: fig = plt.figure(1, figsize=(NUM_CATEGORIES, NUM_CATEGORIES))
grid = ImageGrid(fig, 111, nrows_ncols=(NUM_CATEGORIES, NUM_CATEGORIES), axes_pad=0.05)
i = 0
for category_id, category in enumerate(CATEGORIES):
    for filepath in train[train['category'] == category]['file'].values[:NUM_CATEGORIES]:
        ax = grid[i]
        img = read_img(filepath, (224, 224))
        ax.imshow(img / 255.)
        ax.axis('off')
        if i % NUM_CATEGORIES == NUM_CATEGORIES - 1:
            ax.text(250, 112, filepath.split('/')[1], verticalalignment='center')
            i += 1
plt.show();
```



Validation split

```
In [17]: np.random.seed(seed=SEED)
rnd = np.random.random(len(train))
train_idx = rnd < 0.8
valid_idx = rnd >= 0.8
ytr = train.loc[train_idx, 'category_id'].values
yv = train.loc[valid_idx, 'category_id'].values
len(ytr), len(yv)
```

Out[17]: (1899, 501)

Extract Xception bottleneck features

```
In [18]: INPUT_SIZE = 299
POOLING = 'avg'
x_train = np.zeros((len(train), INPUT_SIZE, INPUT_SIZE, 3), dtype='float32')
for i, file in tqdm(enumerate(train['file'])):
    img = read_img(file, (INPUT_SIZE, INPUT_SIZE))
    x = xception.preprocess_input(np.expand_dims(img.copy(), axis=0))
    x_train[i] = x
print('Train Images shape: {} size: {:,}'.format(x_train.shape, x_train.size))
```

2400it [01:02, 38.27it/s]

Train Images shape: (2400, 299, 299, 3) size: 643,687,200

```
In [19]: Xtr = x_train[train_idx]
Xv = x_train[valid_idx]
print((Xtr.shape, Xv.shape, ytr.shape, yv.shape))
xception_bottleneck = xception.Xception(weights='imagenet', include_top=False, pooling=POOLING)
train_x_bf = xception_bottleneck.predict(Xtr, batch_size=32, verbose=1)
valid_x_bf = xception_bottleneck.predict(Xv, batch_size=32, verbose=1)
print('Xception train bottleneck features shape: {} size: {:,}'.format(train_x_bf.shape, train_x_bf.size))
print('Xception valid bottleneck features shape: {} size: {:,}'.format(valid_x_bf.shape, valid_x_bf.size))
```

((1899, 299, 299, 3), (501, 299, 299, 3), (1899,), (501,))
60/60 [=====] - 13s 141ms/step
16/16 [=====] - 2s 159ms/step
Xception train bottleneck features shape: (1899, 2048) size: 3,889,152
Xception valid bottleneck features shape: (501, 2048) size: 1,026,048

LogReg on Xception bottleneck features

```
In [20]: logreg = LogisticRegression(multi_class='multinomial', solver='lbfgs', random_state=SEED)
logreg.fit(train_x_bf, ytr)
valid_probs = logreg.predict_proba(valid_x_bf)
valid_preds = logreg.predict(valid_x_bf)
```

```
In [21]: print('Validation Xception Accuracy {}'.format(accuracy_score(yv, valid_preds)))
```

Validation Xception Accuracy 0.8363273453093812

Confusion matrix


```
In [22]: cnf_matrix = confusion_matrix(yv, valid_preds)
```

```
In [23]: abbreviation = ['BG', 'Ch', 'CI', 'CC', 'CW', 'FH', 'LSB', 'M', 'SM', 'SP', 'SFC', 'SB']
pd.DataFrame({'class': CATEGORIES, 'abbreviation': abbreviation})
```

```
Out[23]:
```

	class	abbreviation
0	Black-grass	BG
1	Charlock	Ch
2	Cleavers	CI
3	Common Chickweed	CC
4	Common wheat	CW
5	Fat Hen	FH
6	Loose Silky-bent	LSB
7	Maize	M
8	Scentless Mayweed	SM
9	Shepherds Purse	SP
10	Small-flowered Cranesbill	SFC
11	Sugar beet	SB

```
In [24]: fig, ax = plt.subplots(1)
ax = sns.heatmap(cnf_matrix, ax=ax, cmap=plt.cm.Greens, annot=True)
ax.set_xticklabels(abbreviation)
ax.set_yticklabels(abbreviation)
plt.title('Confusion Matrix')
plt.ylabel('True class')
plt.xlabel('Predicted class')
fig.savefig('Confusion matrix.png', dpi=300)
plt.show();
```



Create submission

```
In [25]: x_test = np.zeros((len(test), INPUT_SIZE, INPUT_SIZE, 3), dtype='float32')
        for i, filepath in tqdm(enumerate(test['filepath'])):
            img = read_img(filepath, (INPUT_SIZE, INPUT_SIZE))
            x = xception.preprocess_input(np.expand_dims(img.copy(), axis=0))
            x_test[i] = x
        print('test Images shape: {} size: {}'.format(x_test.shape, x_test.size))

794it [00:08, 90.71it/s]
test Images shape: (794, 299, 299, 3) size: 212,953,182

In [26]: test_x_bf = xception_bottleneck.predict(x_test, batch_size=32, verbose=1)
        print('Xception test bottleneck features shape: {} size: {}'.format(test_x_bf.shape, test_x_bf.size))
        test_preds = logreg.predict(test_x_bf)

25/25 [=====] - 4s 154ms/step
Xception test bottleneck features shape: (794, 2048) size: 1,626,112

In [27]: test['category_id'] = test_preds
        test['species'] = [CATEGORIES[c] for c in test_preds]
        test[['file', 'species']].to_csv('xception_result_v4', index=False)
```

Q3) Screenshots of the result published in Kaggle **Plant Seedlings Classification**

1) Initial Convolutional Neural Networks(CNN) model accuracy was, **0.81863**

Home

Competitions

Datasets

Code

Discussions

Courses

More

Recently Viewed

Plant Seedlings Classif...

Give Me Some Credit

Seedling Classification...

notebook9dd2 added12c

Reading Dataset in Ka...

View Active Events

Search

Plant Seedlings Classification

Determine the species of a seedling from an image

Kaggle

833 teams · 3 years ago

Overview

Data

Code

Discussion

Leaderboard

Rules

Team

My Submissions

Late Submission

Your most recent submission

Name	Submitted	Wait time	Execution time	Score
result_v2.csv	just now	1 seconds	0 seconds	0.81863
Complete				
Jump to your position on the leaderboard				

>_

kaggle competitions submit -c plant-seedlings-classification -f submission.csv -m "Message"


?

Make a submission for Randika Senanayake

Step 1

Upload submission file

2) By using the already trained CNN (Xception model), able to increase accuracy to **0.82619**



Home

Competitions

Datasets


Code


Discussions


Courses


More


Recently Viewed

 Plant Seedlings Classif...

 Image Classification w...

 Seedlings - Pretrained ...

 Seedling Classification...


 VGG08

Search

Playground Prediction Competition

Plant Seedlings Classification

Determine the species of a seedling from an image

 Kaggle · 833 teams · 3 years ago

Overview

Data

Code

Discussion

Leaderboard

Rules

Team

My Submissions

Late Submission

...

Your most recent submission



Name	Submitted	Wait time	Execution time	Score
xception_result_v4.csv	just now	1 seconds	0 seconds	0.82619

Complete

[Jump to your position on the leaderboard](#)

>_

kaggle competitions submit -c plant-seedlings-classification -f submission.csv -m "Message"

Make a submission for [Randika Senanayake](#)